

Leaning Over to Reach Far-Away Objects

Anonymous Authors

Abstract—We study supported reaching to extend a humanoid robot’s workspace by bracing against the environment for stabilization using exclusively local sensing. While current solutions for whole body humanoid control aim at avoiding contact, contact is needed to reach objects that are not directly placed at the edge of a table. We employ a sampling-based model-predictive control (MPC) framework for whole-body leaning, bracing, reaching, and recovery on a Unitree H1-2 humanoid. The controller runs in real-time at 33Hz, transfers sim-to-real, extends reach by more than 40 cm beyond the unbraced standing workspace, and sustains the loaded posture for over four minutes, before pose estimation drifts. The controller reaches the commanded target in 89% of trials and completes the full reach-and-recovery cycle in 81%, across 54 trials spanning six targets and three table heights. We further characterize stability, contact loading, and end-effector precision, and analyze failure modes and sampling-based planner sensitivity. Code, videos, and full hardware details are available at <https://h12leaner.github.io/>

I. INTRODUCTION

Humanoid robots have made significant progress in perception, manipulation, and locomotion, marked by impressive demonstrations of high-speed, dynamic, and contact-rich behaviors inspired by human athletics [1, 2, 3, 4]. However, capability is not defined exclusively by speed. Among humans, physical performance also depends on accurate whole-body coordination, often under sustained load. More broadly, motor-control theory has long characterized skilled movement as the coordination of many biomechanical degrees of freedom to perform physically-constrained tasks [5, 6]. A similarly complete view of physically capable humanoid behavior should thus include not only explosive motion, but also motion that can establish, sustain, and release load-bearing contacts while maintaining balance and task-space precision.

One representative task is reaching objects close to the center of a table or other containers. For example, in industrial battery disassembly [7, 8], many parts are not accessible without leaning into the battery. Humans achieve this by shifting their balance and loading onto available environmental supports. In humanoid robotics, this is challenging as walking controllers are designed to keep the upper body upright. In this work, we study supported reaching: using the upper body to brace against the environment, extend the whole-body workspace beyond the standing limit, and subsequently recovering without locomotion.

We present a sampling-based, contact-implicit model predictive control (CI-MPC) framework for whole-body leaning, bracing and reaching, and recovery on a Unitree H1-2 humanoid. The resulting controller transfers sim-to-real, extends the robot’s reachable workspace by more than 40cm beyond unbraced standing, sustains the loaded posture for more than four minutes, and completes 73% of maximum-extent reaches




Fig. 1: Four stages of leaning and reaching Unitree H1-2 in real (top) and simulation (MuJoCo, bottom). With sampling-based model predictive control, the robot can hold the 1.3 m reaching stage (frame 4), 40 cm past the unbraced standing reach of ≈ 80 cm for over four minutes with EE accuracy of 2.7 ± 0.9 cm.

and 89% of commanded reaches across six targets and three table heights. We further study the whole-body stability, target-tracking error, and steady-state jitter of the controller, as well as the component-level contribution of planner costs and four sampling-based planners.

II. RELATED WORK

a) Whole-body and multi-contact control: Humanoid whole-body control becomes particularly challenging when employing environmental contacts beyond the feet, a problem commonly studied as multi-contact planning and control (MCPC), where a controller must reason over both the robot’s whole-body motion and a set or sequence of environmental support contacts [9]. Common approaches separate these problems into two: first planning a contact sequence, and then optimizing whole-body control trajectories (and contact forces) under a priori contact constraints.

Several recent systems demonstrate the value of upper-body support contacts for humanoid manipulation. Among “classical” control methods, most relevant is SEIKO [10], which performs real-time multi-contact whole-body control on the position-controlled TALOS humanoid, including far-reaching motions that exploit explicitly planned environmental bracing contacts for support. Related work extends this controller with learned multi-support behaviors and automatic support placement for manipulation [11], while other approaches adapt contact locations or whole-body posture to improve stability during multi-contact teleoperation [12]. Among learning-based methods which instead encode whole-body behavior into reusable policies or skill spaces, most relevant to our work are: Opt2Skill [13], which learns to track a dataset of

dynamically feasible supported reaches generated by trajectory optimization with explicitly-planned contacts (Crocodyl [14]), and R²S² [15], which learns a whole-body skill space for large-workspace humanoid reaching, without environmental contacts.

These prior works have demonstrated large-workspace and multi-contact humanoid reaching with learned and optimized methods. We build upon this work with an online whole-body controller that is directly parameterized by geometry and object target commands. We additionally characterize whole-body stability and target-tracking accuracy, commandable workspace, and underlying optimizer operational envelope.

b) MuJoCo MPC: Model predictive control is a common approach for whole-body control of legged robots. MuJoCo MPC (MJPC) [16] is a GUI-based platform interfacing physical simulation in MuJoCo, many classes of MPC planners (sampling and gradient based), and a large set of tasks and robotic embodiments. We therefore employ MJPC for ease of rapid development and iteration. The framework uses iterative linear-quadratic regulator (iLQG) methods with finite-difference derivatives, enabling real-time operation at 20-50Hz on modern hardware. Recent work has demonstrated MJPC with iterative Linear-Quadratic-Regulator (iLQR) planning on the Unitree H1 humanoid [17, 18], showing the promise of MJPC for real-time, real-world locomotion and manipulation on modern computing platforms. Our work extends MJPC deployment to sampling-based planning on the Unitree H1-2 (two more DoFs) **without** requiring a motion-capture system.

III. METHODS

A. Model Predictive Control

We employ MPC within the MuJoCo physics simulator, specifically leveraging MJPC. At each control timestep, the first control u_0 is applied and the planner re-solves a trajectory optimization problem from the updated measured state over a finite-horizon T of 1s:

$$\min_{u_{0:T}} \sum_{t=0}^T c(x_t, u_t) \quad (1)$$

subject to dynamics $x_{t+1} = f(x_t, u_t, \Delta t)$, where $x_t \in \mathbb{R}^{n_x}$ is the state (positions and velocities), $u_t \in \mathbb{R}^{n_u}$ is the control (joint torques or position targets), and $c(\cdot)$ is the running cost over the finite horizon.

MJPC decomposes the cost $c(x, u) = p(l(x, u); R)$, where the base cost is defined as:

$$l(x, u) = \sum_{i=1}^M w_i \cdot n_i(r_i(x, u)) \quad (2)$$

This base cost is a sum of M terms, each comprising a tunable weight $w \in \mathbb{R}_+$, a twice-differentiable norm function $n(\cdot)$, and residual terms $r \in \mathbb{R}^p$. Then to modulate risk sensitivity, an exponential scalar transformation p is applied for some scalar parameter $R \in \mathbb{R}$, where $R > 0$ is risk-averse and $R < 0$ is risk-seeking.

Algorithm 1 One planning iteration of sampling-based M(J)PC

Input: control θ , state x_0 , rollouts n , horizon T

- 1: $\bar{\theta} \leftarrow \text{CENTER}(\theta)$ $\triangleright$ what to sample around
- 2: **for** $k = 1 \dots n$ **do**
- 3: $\epsilon^{(k)} \sim \mathcal{N}(0, \Sigma^{(k)})$ $\triangleright$ SPREAD
- 4: $\theta^{(k)} \leftarrow \bar{\theta} + \epsilon^{(k)}$
- 5: $u_t \leftarrow \text{spline}(\theta^{(k)}; t)$, roll out f from x_0
- 6: $J^{(k)} \leftarrow \frac{1}{T} \sum_{t=0}^T c(x_t, u_t)$ $\triangleright$ Eq. (1)
- 7: **end for**
- 8: $\theta \leftarrow \text{FOLD}(\{\theta^{(k)}\}, \{J^{(k)}\})$ $\triangleright$ how the winner is kept
- 9: apply $\text{spline}(\theta; 0)$; shift θ forward one step

Table I. How planners instantiate Alg. 1. E is the elite set; $k^* = \arg \min_k J^{(k)}$; $w_k \propto e^{-J^{(k)}/\lambda}$. All ablations use $N = 20$ rollouts.

Planner	CENTER	SPREAD	FOLD	N
Predictive sampling (PS) [16]	θ	$\sigma^2 I$	$\theta^{(k^*)}$	20
Model predictive path integral (MPPI) [19]	θ	$\sigma^2 I$	$\sum_k w_k \theta^{(k)}$	20
Cross-entropy method (CEM) [20]	μ_E	Σ_E	μ_E	20
Improved CEM (iCEM) [†] [21]	μ_E	$\Sigma_E, \text{AR}(1)$	μ_E	20

[†]iCEM additionally reuses elites between planning iterations.

Each residual r_i captures a specific objective (e.g., reaching a target, maintaining balance, minimizing control effort). The weights w_i are tuned to trade off competing objectives. MJPC implements this through user-defined sensor (residual) specifications in XML with corresponding weight parameters.

MuJoCo models contact and friction with a convex approximation of underlying non-convex, non-smooth contact dynamics. In lieu of analytical derivatives of the soft contact model over a given timestep $x_{t+1} = f(x_t, u_t)$, MuJoCo approximates model derivatives (dynamics Jacobians) via finite differences:

$$\frac{\partial f}{\partial u} \approx \frac{f(x, u + \epsilon) - f(x, u)}{\epsilon} \quad (3)$$

A similar expression can be constructed for $\partial f / \partial x$ and finite-difference step ϵ [22, 18]. These state and control Jacobians, and their second-order derivative Hessians, are commonly used for planning in MPC, such as iterative Linear Quadratic Regulator (iLQR) or differential dynamic programming (DDP), but often fail to find local minima (e.g., to discover contact) in CI-MPC problems due to un-differentiable contact discontinuities.

Thus, we primarily explore gradient-free sampling-based planners to solve the CI-MPC task of brace contact selection and leaning. At each replanning step it draws a batch of N candidate control sequences $\{u_{0:T}^{(i)}\}_{i=1}^N$, each a perturbation of a nominal spline, rolls every candidate out through the MuJoCo model in parallel, and scores the accumulated cost $\sum_t \ell(x_t^{(i)}, u_t^{(i)})$ along each rollout. Alg. 1 describes one planning iteration of sampling-based MPC, and Tab. I situates four sampling-based planners in Alg. 1. In short: predictive sampling keeps the single best rollout [16], model-predictive path integral (MPPI) forms a cost-weighted average [19], and cross-entropy method (CEM) and its improved version with

TABLE II: Sampling-controller cost terms, grouped by role. The deployed task defines 39 residuals, grouped here into nine functional families. $[x]_+ = \max(0, x)$; d_i^{seat} and d_i^{keep} are clearances to the usable and physical table surfaces; $\mathcal{B}$ is the selected brace set.

Group	Representative residual
Brace	$\min(\bar{d}, [d_i^{\text{seat}} - d_{\text{sat}}]_+)$, $i \in \mathcal{B}$
Keep-out	$[d_{\text{safe}} - d_i^{\text{keep}}]_+$, $i \notin \mathcal{B}$
Trunk clear	$[d_{\text{safe}} - d_{\text{trunk}}]_+$
Reach	$p_{\text{hand}} - p^*$
Balance	$ c_{\text{cp}} - \Pi_{\mathcal{S}}(c_{\text{cp}}) $
Feet & stance	foot rise/tilt, capture, knees, hip posture
Posture	$q - q_{\text{home}}$: pelvis / torso / waist orientation
Height	$[h_{\text{min}} - (z_{\text{head}} - \frac{z_L + z_R}{2})]_+$
Regularization	$\dot{q}$; $u - q_{\text{home}}$

colored noise (iCEM [21]) average a selected elite set and refit the sampling distribution [20]. We utilize CEM as the default planner in following experiments and further study the parameters of these planners to understand component-level effects on the task and sim-to-real mismatch in Sec. IV-E.

Alternative approaches may use analytical derivatives of rigid-body models (Pinocchio [23]) for gradient-based trajectory optimization, such as Crocoddyl [14], which provides differential dynamic programming (DDP) solvers. While such methods offer fast convergence through exact gradients, they require more careful contact model specification and may be less robust to model mismatch than gradient-free, sampling-based approaches.

B. Task Formulation

The task consists of three sequential modes: **(1)** standing, **(2)** reaching toward an object, and **(3)** returning to standing.

Tab. II describes the cost terms used for CI-MPC. Here, $c_{\text{cp}} = c_{\text{CoM}} + 0.3 \dot{c}_{\text{CoM}}$ is the capture point, and $\mathcal{S}$ is the rectangular support region defined by the feet. Once a selected brace is within 10 mm of the table, it is treated as seated for purposes of extending the support region $\mathcal{S}$. The quantities d_i^{seat} and d_i^{keep} denote geometric clearances from the usable bracing surface and the physical table surface, respectively; the former uses a 60 mm inset from the table edge to prevent satisfying the seating objective by contacting the lip. The set $\mathcal{B}$ contains the links selected for bracing by nonzero brace weights. Thus, brace terms encourage those links to approach and seat on the table, while keep-out terms hold all unselected candidate links away from it. While this cost shape is a gray area between explicit and implicit contact planning, MJPC does not enforce a specific contact schedule and nominally requires contact discovery emergent from sampling over rolled-out dynamics.

The deployed configuration (20 rollouts over a 1 s planning horizon) plans at a median 34Hz on a laptop with an AMD Ryzen AI 9 HX 370 CPU. Preliminary testing on a desktop with an Intel i7-14700F CPU showed ≈ 22 Hz planning and increased failure.

C. State Estimation

To estimate the floating-base position and velocity $x_b, v_b \in \mathbb{R}^3$, we fuse proprioceptive measurements, IMU orientation R_b and angular velocity ω_b , together with joint positions and velocities $q, \dot{q}$ with exteroceptive position measurements from LiDAR and AprilTag fiducials. The proprioceptive estimator uses a random-walk Kalman model for base velocity,

$$v_{b,k} = v_{b,k-1} + w_k, \quad z_k = v_{b,k} + \eta_k, \quad (4)$$

where w_k and η_k denote process and measurement noise, respectively.

During standing, each planted foot provides a kinematic observation of base velocity [24]. Assuming zero world-frame velocity of contact point k ,

$$z_k = -[\omega_b \times (p_k(q) - p_b)z + R_b J_k(q) \dot{q}], \quad (5)$$

where J_k is the translational contact Jacobian. Observations from the two feet are incorporated with standard Kalman updates using contact-weighted, outlier-gated measurement noise [25]. Integrating the resulting velocity estimate provides base position, but residual orientation and kinematic errors lead to long-horizon drift [24, 26]. We therefore additionally fuse LiDAR- and AprilTag-based position observations to provide an external reference and bound this accumulated position error.

IV. EXPERIMENTS

We evaluate our methods on a Unitree H1-2 in simulation (MuJoCo) and in the real-world on an adjustable table. Planning and state estimation is done on a laptop with a Ryzen AI 9 HX 370 CPU with the H1-2 on-board computer publishing camera telemetry. We apply a 1 mm thick rubber tape on the table to increase the static coefficient of friction with the H1-2 aluminum body from $\mu \approx 0.4$ to $\mu \approx 0.8$.

A. Estimation

We evaluate three state estimation methods:

1. **Internal:** IMU + leg-kinematic odometry with kinematic-chain guards,
2. **Internal + LiDAR:** the internal estimator fused with FAST-LIO [27],
3. **Internal + Fiducials:** the internal estimator fused with position observations from an external bundle of four AprilTags (36h11, 5–10 cm wide), viewed by the head-mounted Intel RealSense D435i camera.

We evaluate each estimator while the controller actively holds a stand posture. During such a stand the base genuinely sways, so reported base translation no longer isolates error. Since the base naturally sways during the stand, base translation alone cannot distinguish estimator drift from actual motion. We instead use a *foot-anchoring* residual: for a planted foot, forward kinematics from the estimated base pose and measured joint angles should produce a stationary world-frame foot position. Its deviation therefore measures pose error without requiring external ground truth. We report the




Fig. 2: Reach, stability, hand jitter, and brace loading in successful runs. Each panel reports two tasks: reaching a commanded target (left pair) and reaching to maximum extent (right pair), in simulation (orange) and on hardware (green). Bars show medians over successful trials; whiskers show the half-ranges. **Reach** is the gripper-tip distance from the ankle midpoint. **Support margin** is the minimum margin to the active support-region boundary; **hand jitter** is high-pass RMS motion above 1 Hz; and **bracing-arm force** is the measured normal load. Simulation and hardware use the same 37 cm ankle-to-edge stance.




Fig. 3: Estimator drift during an active stand. While the robot actively holds a stand, Z-axis drift is 0.3mm/min, X-axis drift of the Internal + AprilTag Fiducial estimator is 4 mm oscillation..

horizontal and vertical components, along with the inter-foot vector, over the planted-stand interval.

The main difference between the three estimators is horizontal world-frame position (Fig. 3). Internal odometry drifts at +267,mm/min, or roughly 0.48,m over a 2,min hold, while Internal + LiDAR drifts at +210,mm/min. The Fiducial estimator reduces this to +4,mm/min, a $\sim 60\times$ reduction.




Fig. 4: Balance-relevant estimator states during standing. Base-attitude jitter, height drift, inter-foot wander, and reported velocity noise are equivalent across all three estimators, which is why all keep the robot stable.

The remaining metrics are similar across all three (Fig. 4): attitude jitter remains below 0.4° in pitch and 0.2° in roll, height drift is 1–2,mm, inter-foot variation is 14–47,mm, and velocity noise is 7–10,mm/s. Internal is comparable to or better than the fused variants for attitude and inter-foot consistency, while fiducials primarily improve absolute position. Vertical drift is negligible (< 0.3 ,mm over the full trial). We therefore use Internal + Fiducial for the remaining experiments.

B. Leaning, Reaching, and Recovering

We evaluate two tasks: reaching to a commanded target pose (targeted reach) and reaching to the maximum extent afforded by the planner (max reach). In simulation, the controller completes $90 \pm 13\%$ (18/20) of targeted-reach and $80 \pm 18\%$ (16/20) of max-reach trials. On hardware, it completes $80 \pm 20\%$ (12/15) of targeted-reach trials (the max-reach task with the target position explicitly specified) and $73 \pm 22\%$ (11/15) of max-reach trials.

The remainder of this section is devoted to stability analysis of successful runs, described in whole by Fig. 2. We investigate failure in greater detail in Sec. IV-C.

a) Reach: Simulation and hardware reach the same distance. Measured from the ankle midpoint to the gripper tip, the targeted reach settles at 99 ± 1 cm in simulation and 101 ± 4 ,cm on hardware, with the target centroid at 102cm. At max reach the tip extends to 123 ± 4 cm in simulation and 130 ± 10 ,cm on hardware, roughly 45cm beyond the ≈ 80 cm the same stance reaches without bracing.

b) Support margin: Simulation and hardware agree on the stability of the braced posture. The wrench-space margin is 13.2 ± 1.0 cm in simulation and 12.7 ± 1.9 cm on hardware at the targeted reach, and 16.0 ± 0.4 versus 11.7 ± 1.6 cm at max reach. Every successful trial in both conditions keeps a positive margin throughout the brace, and the max-reach hardware margin is the smallest because the robot recovers from the longest lever arm.

c) Bracing-arm force: The simulated controller loads the forearm harder than the real one. Median brace load is 182 ± 31 ,N at the targeted reach and 174 ± 50 N at max reach in simulation, against 121 ± 47 and 124 ± 30 N on hardware. Two factors explain the gap: the simulated arm is unclamped,




Fig. 5: CoM support margin and reconstructed brace-contact force over hardware trials. Time is normalized to the commanded phase boundaries for each trial. Top: minimum distance from the horizontal CoM projection to the **foot-only support polygon**; negative values lie outside the foot support region. Bottom: normal contact force estimated quasi-statically from measured bracing-arm joint torques through the contact Jacobian transpose. The reconstruction gives 4.3 N RMSE against ground-truth force in simulation.




Fig. 6: Gripper tip tracking. Wrist-camera measured tip error relative to the aim point through the 5 s hold, mean $\pm$ 1 sd over 47 hardware holds across six targets and three table heights. Depth carries a small bias as the hand drives in along the approach (-1.4 ± 1.3 cm), while the grasp-critical lateral and height axes stay well centred at $+0.2 \pm 1.3$ and $+1.1 \pm 1.3$ cm within the jaw aperture.

whereas the deployed arm saturates near 100N of shoulder-pitch torque capacity, and the simulated robot leans deeper into the table, which also raises its support margin. Hardware force is reconstructed quasi-statically from bracing-arm joint torques (Fig. 5).

d) Hand jitter: High-frequency hand motion during the hold matches across the sim-to-real gap. Jitter above 1Hz is 1.5 ± 0.4 mm RMS in simulation and 1.5 ± 0.4 mm on hardware at the targeted reach. At max reach hardware jitter rises to 4.9 ± 2.7 mm against 1.3 ± 0.2 mm in simulation, consistent with the smaller support margin at that posture.

1) Target Position Accuracy: To quantify end effector precision, we command the gripper tip to a fixed Cartesian target

in the table frame (0.55 m past the near edge, 0.16 m above the surface) and evaluate placement accuracy at convergence. Across 10 autonomous trials, the tip converges to 2.8 ± 1.3 cm in the plane transverse to the reach axis, shown in Fig. 6, with a lateral component of $+1.3 \pm 2.1$ cm and a height component of -0.5 ± 1.8 cm.

C. Failure Analysis

Successful and failed trials separate primarily in how they load and maintain the brace. We analyze the 15 max-reach hardware trials (11 successes, 4 failures). During bracing, both groups move the CoM beyond the foot-only support polygon (Fig. 5). Successful trials gradually recover to a positive margin up through brace release. In failed trials, the robot slowly falls forward after the initial brace, reaching approximately -0.14 m overall and below -0.10 m relative to the front edge before triggering an emergency stop.

The corresponding contact-force profiles further distinguish successful and failing loading strategies (Fig. 5). Successful trials load the forearm progressively, reaching approximately 125 N on average, with per-trial peaks of 151 ± 24 N, and unload gradually during recovery. Failed trials instead develop approximately 150 N near brace onset and subsequently maintain a more variable 100–150 N load until collapse.

Shoulder torque and position measurements further diagnose failure. Successful trials sustain approximately $+10$ N m of shoulder-pitch torque while braced, whereas failures reach -20 to -25 N m, reflecting configuration ($+0.14 \pm 0.08$ versus -0.45 ± 0.10 rad). To provide some physical intuition, this difference reflects the initial brace placing the shoulder link behind (success) or ahead (failure) of the elbow, and currently is best attributed to sampling variance in the current planner configuration.

D. Table Height and Target Parameterization

We demonstrate that the proposed controller can reach not just one optimized target, but also be commanded by the geometry of the scene. We vary both the table height and reach target while otherwise leaving the controller unchanged and report 89% (48/54) reach and 81% (44/54) full recovery success across three table heights, 0.935, 0.985, and 1.035 m (± 5 cm), and six targets arranged across a 20 cm by 10 cm grid, as shown in Fig. 7. The target positions are set to, along the forward X-axis, 92 cm (7 cm inset from the maximum reach) and 82 cm, and ± 10 cm along the Y-axis from the original max and targeted reach position. For each table height, the measured table surface is updated in the planner model and the target position is then provided directly to the same controller; otherwise, no cost weights or control parameters are changed between conditions. The target is a 5 cm cube observed through AprilTags by the wrist-mounted camera (Fig. 7, left), which continuously updates the commanded target (visual-servoing). For each individual target, we run three trials consecutively without resetting the robot, with no quantifiable degradation in stability or task performance, for 18 reaches per table height.

Lean and Reach Across Multiple Table Heights (n=3) and Targets (n=6) (XY-View)




Fig. 7: Commanded reaching across table heights and targets. The same controller is evaluated at three table heights and six target positions, with three trials per condition. Dots show wrist-camera measured hold positions in the table plane; crosses denote target positions, dotted circles show a 4 cm radius, and $\times$ marks conditions where the brace failed before reaching. All targets are completed at the lower and nominal table heights; at the tallest table, the two deeper lateral targets exceed the current bracing workspace.




Fig. 8: Brace-cost ablation. Brace-hold success rate versus the number of active brace cost terms (Wilson 95% intervals; per-point n annotated).

The robot completes all 18/18 holds at 0.935 and 0.985 m, and 12/18 at 1.035 m. The two deepest lateral targets at 1.035 m fail during the initial lean, where the forearm brace collapses before the reach is established. Then, at 0.935 m, we report four failures to recover from the reaching position, similarly due to the forearm brace becoming leveraged too far forward. Across the 48 completed holds, the wrist-camera measurement gives a mean 3D target error of 2.7 ± 0.9 cm, with 44/48 holds within 4 cm of the commanded aim point. Mean errors remain similar across table heights: 2.8, 2.5, and 3.0 cm, respectively. Figure 7 shows the corresponding top-down hold locations.

These results suggest that supported reaching as proposed can be used as a manipulation primitive: perception can provide a support height and object pose, and the whole-body controller can realize the corresponding reach without a task-specific motion or controller retuning. We note that


planning and modeling improvements, particularly brace pose and contact selection, can drastically improve performance.

E. Cost and Planner Ablation


a) *Brace Cost Ablation:* The intended brace is shaped by several cost terms (the *Brace* group in Table II). Figure 8 ablates them: brace-hold success (a stable, loaded forearm seat) rises monotonically with the number of active brace terms, from 56% (24/43) with six terms to 100% (15/15) with the full nine. Panel (b) reports the per-term contribution and contrasts the final-config brace success (100%, 15/15) with the full recovery-pipeline success (73%, 11/15).

b) *Planner Ablation:* One may be curious as to why the proposed controller can be effective while planning at a relatively slow 33 Hz, especially since recent whole-body, MuJoCo-based MPC systems for legged robots emphasize high-rate control (50 Hz iLQR with 300 Hz smoothing [18], or 100 Hz MPPI [28]). To understand the operating limits of our method, we sweep plan rate, latency, and planner-plant mismatch in simulation. Broadly, we find that supported reaching is a slowly varying posture stabilized by the joint PD controller; planning demand is concentrated around changes in support and planner choice matters most as the model deviates from the plant. We note that in preliminary testing, gradient-based methods like iLQR, iLQG, and DDP failed completely with the controller as currently constructed.

For a common 0.01 rad sampling standard deviation and a 1 s control trajectory with piecewise-constant, zero-order interpolation of the three control knots (e.g., each sampled action “knot” or point $u_i \dots u_K$ in the N -order spline representation of the trajectory is applied, unaltered), all four planners solve the nominal task at 33 Hz planning (4–6/6). How sampled actions evolve over the 1 s rollout strongly affects success: zero-order interpolation holds actions fixed and preserves their effect, whereas linear or cubic interpolation



(a) **Configuration sensitivity.** Completion over sampling standard deviation and plan rate for CEM, predictive sampling (PS), and MPPI using the hold representation. Bottom: CEM elite count k , MPPI temperature λ , and rollout count N at 33 plans/s. CEM retains the widest successful noise-rate basin; each cell uses six seeds and outlined cells complete at least 4/6.



(b) **Model mismatch.** Completion at 33 plans/s when plant mass, joint stiffness, or sliding friction differs from the planner model. Bottom: effect of increasing sampling standard deviation at nominal mass (hollow) and $1.10\times$ mass (filled). CEM is most robust overall, while mass error is difficult for every planner. Error bars are Wilson 95% intervals over 12 seeds.

Fig. 9: Sampling-based planner ablation: sensitivity to configuration and to model error. Hold spline, $\sigma = 0.01$ rad, $N = 20$, $k = 6$ and $\lambda = 1$ unless varied.

immediately blends toward later samples. Across interpolation order, knot count (trajectory resolution), and noise correlation, configurations whose initial perturbation remains correlated with the executed action for $\gtrsim 0.22$ s complete 5–6/6 trials, while those decorrelating by 0.20 s complete only 0–2/6. iCEM succeeds with cubic interpolation because its temporally correlated sampling noise preserves the candidate action over this interval. In short, the sampled action must persist long

enough for its dynamic effect (contact) to occur in the rollout.

With the trajectory representation fixed, the planners differ through the FOLD update in Alg. 1: PS keeps the minimum-cost rollout, CEM averages the lowest-cost elites and refits their covariance, and MPPI forms a softmax-weighted average. These update rules produce different operating ranges as sampling noise and plan rate vary (Fig. 9a). CEM succeeds in 16/20 tested σ -rate conditions, compared with 12/20 for PS and 10/20 for MPPI. Then, at the nominal $\sigma = 0.01$ rad, all planners fail at 10 Hz planning rate, although standing remains stable even at 5 plans/s. For CEM, increasing the sampling scale partially recovers 10 Hz operation (4/6). Latency is less restrictive: at 33 Hz, 100 ms of uncompensated delay retains 5–6/6 success because new plans still arrive every 30 ms.

Finally, we introduce controlled planner-plant mismatch in mass, joint stiffness, and sliding friction as a proxy for parametric components of the sim-to-real gap (Fig. 9b). CEM completes 80/108 trials, compared with 62/108 for iCEM, 47/108 for PS, and 42/108 for MPPI. CEM is especially tolerant to stiffness and friction error, while mass mismatch degrades every planner. Increasing sampling variance also improves CEM and iCEM tolerance to mass error with little nominal penalty, whereas PS and MPPI lose nominal successes. Elite averaging therefore gives CEM a wider operating range over both plan rate and parametric model error.

V. DISCUSSION

We demonstrate whole-body multiple contact planning and control, e.g. leaning over and reaching far-away objects, with a relatively simple sampling-based MPC framework on a Unitree H1-2 humanoid, sim-to-real. We characterize how our proposed method discovers, loads, sustains, and releases an environmental support contact on a H1-2 without an explicit contact schedule or external motion-capture technology.

Future work includes further study of brace contact planning, extending the framework to object retrieval for varying payload mass while under a brace, generalizing across support surfaces, inclinations and heights, deploying humanoid skin to improve contact dynamics and contact sensing resolution, and exploring GPU-parallelizable and/or learned methods to improve planning rate and quality. Of note—we observe that a MacBook M5 Pro chip and Intel Ultra 7 265KF CPU can plan at speeds of 60 Hz and 45 Hz, respectively, on the lean and reach task. Silicon improvements, together with planning and control optimizations, may further facilitate online MPC for such MCPC problems.

ACKNOWLEDGMENTS

AI Usage statement.

REFERENCES

- [1] L. Yang, X. Huang, Z. Wu, A. Kanazawa, P. Abbeel, C. Sferazza, C. K. Liu, R. Duan, and G. Shi, “Omni-target: Interaction-preserving data generation for humanoid whole-body loco-manipulation and scene interaction,” 2026. [Online]. Available: <https://arxiv.org/abs/2509.26633>
- [2] Y. Ma, A. Cramariuc, F. Farshidian, and M. Hutter, “Learning coordinated badminton skills for legged manipulators,” *Science*

- Robotics*, vol. 10, no. 102, p. eadu3922, 2025. [Online]. Available: <https://www.science.org/doi/abs/10.1126/scirobotics.adu3922>
- [3] Z. Zhuang, S. Yao, and H. Zhao, "Humanoid parkour learning," 2024. [Online]. Available: <https://arxiv.org/abs/2406.10759>
 - [4] Z. Zhang, H. Lu, Y. Lian, Z. Chen, Y. Liu, C. Lin, H. Xue, Z. Zeng, Z. Qi, S. Zheng, Q. Luan, J. Wang, J. Xing, H. Wang, and L. Yi, "Learning athletic humanoid tennis skills from imperfect human motion data," 2026. [Online]. Available: <https://arxiv.org/abs/2603.12686>
 - [5] N. A. Bernstein, *The Co-ordination and Regulation of Movements*. Oxford: Pergamon Press, 1967.
 - [6] E. Todorov and M. I. Jordan, "Optimal feedback control as a theory of motor coordination," *Nature Neuroscience*, vol. 5, no. 11, pp. 1226–1235, 2002.
 - [7] Z. Allen, M. Conway, L. Antieau, A. Ponraj, and N. Correll, "Robotic agentic platform for intelligent electric vehicle disassembly," 2026. [Online]. Available: <https://arxiv.org/abs/2603.18520>
 - [8] M. Conway, W. Xie, A. Devaraj, Y. Zhang, N. Pudasaini, M. Feit, A. Abid, Z. Allen, C. Liu, X. Tan, J. Lavering, J. Chen, L. Antieau, A. V. Pischke, A. Roncone, Z. Sunberg, and N. Correll, "Golem: Modular humanoid autonomy towards electric vehicle battery disassembly," 2026. [Online]. Available: <https://arxiv.org/abs/2608.21550>
 - [9] P. Ferrari, L. Rossini, F. Ruscelli, A. Laurenzi, G. Oriolo, N. G. Tsagarakis, and E. Mingo Hoffman, "Multi-contact planning and control for humanoid robots: Design and validation of a complete framework," *Robotics and Autonomous Systems*, vol. 166, p. 104448, 2023.
 - [10] Q. Rouxel, S. Ivaldi, and J.-B. Mouret, "Multi-contact whole-body force control for position-controlled robots," *IEEE Robotics and Automation Letters*, vol. 9, no. 6, pp. 5639–5646, 2024.
 - [11] Q. Rouxel, A. Ferrari, S. Ivaldi, and J.-B. Mouret, "Flow matching imitation learning for multi-support manipulation," in *2024 IEEE-RAS 23rd International Conference on Humanoid Robots (Humanoids)*. IEEE, 2024, pp. 528–535.
 - [12] S. McCrory, R. Orsolino, D. Thanki, L. Penco, and R. J. Griffin, "Stability-aware retargeting for humanoid multi-contact teleoperation," *IEEE Robotics and Automation Letters*, vol. 10, no. 12, pp. 12397–12404, 2025.
 - [13] F. Liu, Z. Gu, Y. Cai, Z. Zhou, H. Jung, J. Jang, S. Zhao, S. Ha, Y. Chen, D. Xu, and Y. Zhao, "Opt2skill: Imitating dynamically-feasible whole-body trajectories for versatile humanoid loco-manipulation," *IEEE Robotics and Automation Letters*, 2025.
 - [14] C. Mastalli, R. Budhiraja, W. Merkt, G. Saurel, B. Hammoud, M. Naveau, J. Carpentier, L. Righetti, S. Vijayakumar, and N. Mansard, "Crocodyl: An Efficient and Versatile Framework for Multi-Contact Optimal Control," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2020.
 - [15] Z. Zhang, C. Chen, H. Xue, J. Wang, S. Liang, Y. Liu, Z. Zhang, H. Wang, and L. Yi, "Unleashing humanoid reaching potential via real-world-ready skill space," *IEEE Robotics and Automation Letters*, vol. 11, no. 2, pp. 2082–2089, 2026.
 - [16] T. Howell, N. Gileadi, S. Tunyasuvunakool, K. Zakka, T. Erez, and Y. Tassa, "Predictive sampling: Real-time behaviour synthesis with mujoco," 2022. [Online]. Available: <https://arxiv.org/abs/2212.00541>
 - [17] M. Meser, A. Bhatt, B. Belousov, and J. Peters, "MuJoCo MPC for humanoid control: Evaluation on HumanoidBench," in *40th Anniversary of the IEEE International Conference on Robotics and Automation (ICRA@40)*. IEEE, 2024. [Online]. Available: <https://arxiv.org/abs/2408.00342>
 - [18] J. Z. Zhang, T. A. Howell, Z. Yi, C. Pan, G. Shi, G. Qu, T. Erez, Y. Tassa, and Z. Manchester, "Whole-body model-predictive control of legged robots with MuJoCo," in *2026 IEEE International Conference on Robotics and Automation (ICRA)*, 2026.
 - [19] G. Williams, N. Wagener, B. Goldfain, P. Drews, J. M. Rehg, B. Boots, and E. A. Theodorou, "Information theoretic MPC for model-based reinforcement learning," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2017, pp. 1714–1721.
 - [20] R. Y. Rubinstein and D. P. Kroese, *The Cross-Entropy Method: A Unified Approach to Combinatorial Optimization, Monte Carlo Simulation and Machine Learning*. New York: Springer, 2004.
 - [21] C. Pinneri, S. Sawant, S. Blaes, J. Achterhold, J. Stueckler, M. Rolinek, and G. Martius, "Sample-efficient cross-entropy method for real-time planning," in *Conference on Robot Learning (CoRL)*, 2020.
 - [22] E. Todorov, T. Erez, and Y. Tassa, "Mujoco: A physics engine for model-based control," in *2012 IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2012, pp. 5026–5033.
 - [23] J. Carpentier, G. Saurel, G. Buondonno, J. Mirabel, F. Lamiraux, O. Stasse, and N. Mansard, "The pinocchio c++ library – a fast and flexible implementation of rigid body dynamics algorithms and their analytical derivatives," in *IEEE International Symposium on System Integrations (SII)*, 2019.
 - [24] M. Bloesch, M. Hutter, M. Hoepflinger, S. Leutenegger, C. Gehring, C. D. Remy, and R. Siegwart, "State estimation for legged robots - consistent fusion of leg kinematics and IMU," in *Proceedings of Robotics: Science and Systems*, Sydney, Australia, July 2012.
 - [25] M. Camurri, M. Fallon, S. Bazeille, A. Radulescu, V. Barasuol, D. G. Caldwell, and C. Semini, "Probabilistic contact estimation and impact detection for state estimation of quadruped robots," *IEEE Robotics and Automation Letters*, vol. 2, no. 2, pp. 1023–1030, 2017.
 - [26] N. Rotella, M. Bloesch, L. Righetti, and S. Schaal, "State estimation for a humanoid robot," in *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, Chicago, IL, USA, 2014, pp. 952–958.
 - [27] W. Xu, Y. Cai, D. He, J. Lin, and F. Zhang, "FAST-LIO2: Fast direct lidar-inertial odometry," *IEEE Transactions on Robotics*, vol. 38, no. 4, pp. 2053–2073, 2022.
 - [28] J. Alvarez-Padilla, J. Z. Zhang, S. Kwok, J. M. Dolan, and Z. Manchester, "Real-time whole-body control of legged robots with model-predictive path integral control," in *2025 IEEE International Conference on Robotics and Automation (ICRA)*, 2025, pp. 14721–14727.